

day-35-40-handeling-missing-value

May 26, 2025

#Note: - If Numaric Column contian <5% missing value and Randomly. Then you can use —>mean,median stratgy for handeling missing value

- but when not randomly use ‘End of Distribution’-1)Q1-1.5IQR 2)Q1+1.5IQR where IQR=Q3-Q1
- When data is Numercal Catagorical and Missing value >10% .Then use Arbitrary Missing value handeling or insert -1 where NaN value.
- When Data String Catagorical and Missing value >10% .Then use “missing/empty” string as a new Catagory to NaN value

```
# si=SimpleImputer(stratgy='constant',fill_value=-1 | "missing/empty")
X_train['new']=si.fit_transfrom(X_train)
```

```
[ ]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
[ ]: df=pd.read_csv('/content/Titanic-Dataset.csv')
df.head(3)
```

```
[ ]: PassengerId  Survived  Pclass  \
0             1         0         3
1             2         1         1
2             3         1         3
```

```
                                Name    Sex  Age  SibSp  \
0                Braund, Mr. Owen Harris  male  22.0     1
1  Cumings, Mrs. John Bradley (Florence Briggs Th...  female  38.0     1
2                Heikkinen, Miss. Laina  female  26.0     0
```

```
   Parch    Ticket   Fare Cabin Embarked
0      0   A/5 21171   7.2500   NaN        S
1      0    PC 17599  71.2833   C85        C
2      0 STON/O2. 3101282   7.9250   NaN        S
```

```
[ ]: df=df.iloc[:,[1,2,4,5,6,7,9,10,11]]
df.head(3)
```

```
[ ]:      Survived  Pclass      Sex   Age  SibSp  Parch      Fare Cabin Embarked
0         0         3    male  22.0     1     0   7.2500   NaN      S
1         1         1  female  38.0     1     0  71.2833   C85      C
2         1         3  female  26.0     0     0   7.9250   NaN      S
```

```
[ ]: df['Cabin'].isnull().sum()/df.shape[0]*100
```

```
[ ]: np.float64(77.10437710437711)
```

```
[ ]: df.drop('Cabin',axis=1,inplace=True)
df.head(3)
```

```
[ ]:      Survived  Pclass      Sex   Age  SibSp  Parch      Fare Embarked
0         0         3    male  22.0     1     0   7.2500      S
1         1         1  female  38.0     1     0  71.2833      C
2         1         3  female  26.0     0     0   7.9250      S
```

```
[ ]: from sklearn.model_selection import train_test_split
from sklearn.impute import SimpleImputer,MissingIndicator
from sklearn.preprocessing import OneHotEncoder,OrdinalEncoder,LabelEncoder
from sklearn.preprocessing import StandardScaler,MinMaxScaler
from sklearn.preprocessing import FunctionTransformer,PowerTransformer
from sklearn.compose import ColumnTransformer,make_column_transformer
from sklearn.pipeline import Pipeline,make_pipeline
import scipy.stats as stats

from sklearn.linear_model import LogisticRegression,LinearRegression
```

#Casually Implement of MissingIndicator

```
[ ]: df.head(2)
```

```
[ ]:      Survived  Pclass      Sex   Age  SibSp  Parch      Fare Embarked
0         0         3    male  22.0     1     0   7.2500      S
1         1         1  female  38.0     1     0  71.2833      C
```

```
[ ]: mi=MissingIndicator()
df['new_age']=mi.fit_transform(df[['Age']])
df.sample(10)
```

```
[ ]:      Survived  Pclass      Sex   Age  SibSp  Parch      Fare Embarked  new_age
246         0         3  female  25.0     0     0   7.7750      S    False
358         1         3  female   NaN     0     0   7.8792      Q     True
473         1         2  female  23.0     0     0  13.7917      C    False
280         0         3    male  65.0     0     0   7.7500      Q    False
807         0         3  female  18.0     0     0   7.7750      S    False
53          1         2  female  29.0     1     0  26.0000      S    False
```

785	0	3	male	25.0	0	0	7.2500	S	False
7	0	3	male	2.0	3	1	21.0750	S	False
279	1	3	female	35.0	1	1	20.2500	S	False
240	0	3	female	NaN	1	0	14.4542	C	True

```
[ ]: df.shape
```

```
[ ]: (891, 9)
```

```
[ ]: X_train,X_test,y_train,y_test=train_test_split(df.  
↳drop(['Survived','new_age'],axis=1),df['Survived'],test_size=0.  
↳2,random_state=42)  
X_train.head(2)
```

```
[ ]:      Pclass  Sex  Age  SibSp  Parch  Fare Embarked  
331      1  male  45.5     0     0  28.5         S  
733      2  male  23.0     0     0  13.0         S
```

```
[ ]: age_pipe=Pipeline([  
    ('imp_age',SimpleImputer(strategy='median'))  
)  
fare_pipe=Pipeline([  
    ('normalize_fare',PowerTransformer(method='yeo-johnson'))  
)  
  
sex_emb_pipe=Pipeline([  
    ('sex_imp',SimpleImputer(strategy='most_frequent')),  
    ('ohe_sex_emb',OneHotEncoder(dtype=np.  
↳int32,drop='first',sparse_output=False,handle_unknown='ignore'))  
)
```

```
[ ]: CT1=make_column_transformer(  
    (age_pipe,['Age']),  
    (fare_pipe,['Fare']),  
    (sex_emb_pipe,['Sex','Embarked']),  
    remainder='passthrough'  
)
```

```
[ ]: CT1.fit_transform(X_train)
```

```
[ ]: array([[45.5      ,  0.47999826,  1.          , ...,  1.          ,  
         0.          ,  0.          ],  
        [23.       , -0.28375264,  1.          , ...,  2.          ,  
         0.          ,  0.          ],  
        [32.       , -0.77244668,  1.          , ...,  3.          ,  
         0.          ,  0.          ],  
        ...,  
        ...])
```

```
[41.      , -0.20313477,  1.      , ...,  3.      ,
   2.      ,  0.      ],
[14.      ,  1.78386853,  0.      , ...,  1.      ,
   1.      ,  2.      ],
[21.      ,  1.40043606,  1.      , ...,  1.      ,
   0.      ,  1.      ]])
```

```
[ ]: pipe_before=make_pipeline(CT1,StandardScaler(),LogisticRegression())

pipe_before.fit(X_train,y_train)
```

```
/usr/local/lib/python3.11/dist-
packages/sklearn/compose/_column_transformer.py:1667: FutureWarning:
The format of the columns of the 'remainder' transformer in
ColumnTransformer.transformers_ will change in version 1.7 to match the format
of the other transformers.
At the moment the remainder columns are stored as indices (of type int). With
the same ColumnTransformer configuration, in the future they will be stored as
column names (of type str).
To use the new behavior now and suppress this warning, use
ColumnTransformer(force_int_remainder_cols=False).
```

```
warnings.warn(
```

```
[ ]: Pipeline(steps=[('columntransformer',
                      ColumnTransformer(remainder='passthrough',
                                         transformers=[('pipeline-1',
                                                         Pipeline(steps=[('imp_age',
                                                                 SimpleImputer(strategy='median'))]),
                                                         ['Age']),
                                                         ('pipeline-2',
                                                         Pipeline(steps=[('normalize_fare',
                                                                 PowerTransformer()))]),
                                                         ['Fare']),
                                                         ('pipeline-3',
                                                         Pipeline(steps=[('sex_imp',
                                                                 SimpleImputer(strategy='most_frequent')),
                                                                 ('ohe_sex_emb',
                                                                 OneHotEncoder(drop='first',
                                                                 dtype=<class 'numpy.int32'>,
                                                                 handle_unknown='ignore',
                                                                 sparse_output=False))]),
                                                         ['Sex', 'Embarked'])])),
                      ('standardscaler', StandardScaler()),
                      ('logisticregression', LogisticRegression())])
```

```
[ ]: y_pred=pipe_before.predict(X_test)

from sklearn.metrics import accuracy_score
accuracy_score(y_pred,y_test)
```

```
[ ]: 0.7988826815642458
```

```
[ ]: from sklearn.model_selection import cross_val_score
cross_val_score(pipe_before,X_train,y_train,cv=5,scoring='accuracy').mean()
```

```
[ ]: np.float64(0.7864572047670639)
```

##after Missing Indicator(builtin and manually)

```
[ ]: age_p=Pipeline([
    # ('mi_age',MissingIndicator()),
    ('imu_age',SimpleImputer(add_indicator=True))
])
```

```
[ ]: CT2=make_column_transformer(
    (age_p,['Age']),
    (fare_pipe,['Fare']),
    (sex_emb_pipe,['Sex','Embarked']),
    remainder='passthrough'
)
CT2.fit_transform(X_train)
```

```
[ ]: array([[45.5      ,  0.      ,  0.47999826, ...,  1.      ,
          0.      ,  0.      ],
 [23.      ,  0.      , -0.28375264, ...,  2.      ,
          0.      ,  0.      ],
 [32.      ,  0.      , -0.77244668, ...,  3.      ,
          0.      ,  0.      ],
 ...,
 [41.      ,  0.      , -0.20313477, ...,  3.      ,
          2.      ,  0.      ],
 [14.      ,  0.      ,  1.78386853, ...,  1.      ,
          1.      ,  2.      ],
 [21.      ,  0.      ,  1.40043606, ...,  1.      ,
          0.      ,  1.      ]])
```

```
[ ]: pipe_after=make_pipeline(CT2,StandardScaler(),LogisticRegression())
pipe_after.fit(X_train,y_train)
```

/usr/local/lib/python3.11/dist-packages/sklearn/compose/_column_transformer.py:1667: FutureWarning:
The format of the columns of the 'remainder' transformer in
ColumnTransformer.transformers_ will change in version 1.7 to match the format

of the other transformers.

At the moment the remainder columns are stored as indices (of type int). With the same ColumnTransformer configuration, in the future they will be stored as column names (of type str).

To use the new behavior now and suppress this warning, use `ColumnTransformer(force_int_remainder_cols=False)`.

```
warnings.warn(
```

```
[ ]: Pipeline(steps=[('columntransformer',
                      ColumnTransformer(remainder='passthrough',
                      transformers=[('pipeline-1',
                                   Pipeline(steps=[('imu_age',
                                                    SimpleImputer(add_indicator=True))),
                                   ['Age']),
                                   ('pipeline-2',
                                   Pipeline(steps=[('normalize_fare',
                                                    PowerTransformer()))],
                                   ['Fare']),
                                   ('pipeline-3',
                                   Pipeline(steps=[('sex_imp',
                                                    SimpleImputer(strategy='most_frequent')),
                                                    ('ohe_sex_emb',
                                                    OneHotEncoder(drop='first',
                                                                    dtype=<class 'numpy.int32'>,
                                                                    handle_unknown='ignore',
                                                                    sparse_output=False))]),
                                   ['Sex', 'Embarked'])])),
                      ('standardscaler', StandardScaler()),
                      ('logisticregression', LogisticRegression())])
```

```
[ ]: y_pred=pipe_after.predict(X_test)
      accuracy_score(y_pred,y_test)
```

```
[ ]: 0.8100558659217877
```

```
[ ]: from sklearn.model_selection import cross_val_score
      cross_val_score(pipe_after,X_train,y_train,cv=5,scoring='accuracy').mean()
```

```
[ ]: np.float64(0.7864572047670639)
```

#check manually

```
[ ]: df.head()
      #True and False jodi bool type hoi tahole ML a Encode korte hobe na
```

```
[ ]:   Survived  Pclass    Sex  Age  SibSp  Parch    Fare  Embarked  new_age
      0         0        3  male  22.0    1      0   7.2500         S      False
```

1	1	1	female	38.0	1	0	71.2833	C	False
2	1	3	female	26.0	0	0	7.9250	S	False
3	1	1	female	35.0	1	0	53.1000	S	False
4	0	3	male	35.0	0	0	8.0500	S	False

```
[ ]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 9 columns):
#   Column      Non-Null Count  Dtype
---  -
0   Survived    891 non-null    int64
1   Pclass      891 non-null    int64
2   Sex         891 non-null    object
3   Age         714 non-null    float64
4   SibSp       891 non-null    int64
5   Parch       891 non-null    int64
6   Fare        891 non-null    float64
7   Embarked    889 non-null    object
8   new_age     891 non-null    bool
dtypes: bool(1), float64(2), int64(4), object(2)
memory usage: 56.7+ KB
```

```
[ ]: X_train,X_test,y_train,y_test=train_test_split(df.
↳drop(['Survived'],axis=1),df['Survived'],test_size=0.2,random_state=42)
X_train.head(2)
```

```
[ ]:      Pclass  Sex  Age  SibSp  Parch  Fare  Embarked  new_age
331      1  male  45.5    0    0  28.5         S      False
733      2  male  23.0    0    0  13.0         S      False
```

```
[ ]: age_pipe2=Pipeline([
    ('age_imp',SimpleImputer())
])
fare_pipe2=Pipeline([
    ('normalization_fare',PowerTransformer())
])

sex_emb_pipe2=Pipeline([
    ('sex_imp',SimpleImputer(strategy='most_frequent')),
    ('ohe_sex_emb',OneHotEncoder(dtype=np.
↳int32,drop='first',sparse_output=False,handle_unknown='ignore'))
])

CT3=make_column_transformer(
    (age_pipe2,['Age']),
```

```

(fare_pipe2,['Fare']),
(sex_emb_pipe2,['Sex','Embarked']),
remainder='passthrough'
)

CT3.fit_transform(X_train)

```

```

[ ]: array([[45.5, 0.47999826231989395, 1, ..., 0, 0, False],
           [23.0, -0.28375263784654847, 1, ..., 0, 0, False],
           [32.0, -0.7724466813890875, 1, ..., 0, 0, False],
           ...,
           [41.0, -0.20313476693661398, 1, ..., 2, 0, False],
           [14.0, 1.7838685328829773, 0, ..., 1, 2, False],
           [21.0, 1.400436058476949, 1, ..., 0, 1, False]], dtype=object)

```

```

[ ]: pipe3=make_pipeline(CT3,StandardScaler(),LogisticRegression())
      pipe3.fit(X_train,y_train)

```

/usr/local/lib/python3.11/dist-packages/sklearn/compose/_column_transformer.py:1667: FutureWarning:
The format of the columns of the 'remainder' transformer in
ColumnTransformer.transformers_ will change in version 1.7 to match the format
of the other transformers.
At the moment the remainder columns are stored as indices (of type int). With
the same ColumnTransformer configuration, in the future they will be stored as
column names (of type str).
To use the new behavior now and suppress this warning, use
ColumnTransformer(force_int_remainder_cols=False).

```
warnings.warn(
```

```

[ ]: Pipeline(steps=[('columntransformer',
                      ColumnTransformer(remainder='passthrough',
                                         transformers=[('pipeline-1',
                                                         Pipeline(steps=[('age_imp',
                                                                                 SimpleImputer()))],
                                                                                 ['Age']),
                                                         ('pipeline-2',
                                                         Pipeline(steps=[('normalization_fare',
                                                                                 PowerTransformer()))],
                                                                                 ['Fare']),
                                                         ('pipeline-3',
                                                         Pipeline(steps=[('sex_imp',
                                                                                 SimpleImputer(strategy='most_frequent')),
                                                                                 ('ohe_sex_emb',
                                                                                 OneHotEncoder(drop='first',
                                                                                 dtype=<class 'numpy.int32'>,

```



```

handle_unknown='ignore',
sparse_output=False))],
                                ['Sex', 'Embarked']]])),
('standardscaler', StandardScaler()),
('logisticregression', LogisticRegression()))

```

```
[ ]: y_pred3=pipe3.predict(X_test)
accuracy_score(y_pred3,y_test)
```

```
[ ]: 0.8100558659217877
```

```
[ ]: cross_val_score(pipe3,X_train,y_train,cv=5,scoring='accuracy').mean()
```

```
[ ]: np.float64(0.7864572047670639)
```

#Grid SearchCV: Find automatically better parameter for any Transformation which increase accuracy [SEE Latter]

```
[ ]: # from sklearn.model_selection import GridSearchCV
# pram_grid={
#     'A':[0.01, 0.1, 1, 10],
#     'imp_age':['mean','median'],
#     'imp_cat':['most_frequent','constant']
# }
# grid_search=GridSearchCV(pipe3,pram_grid,cv=5,scoring='accuracy')
# grid_search.fit(X_train,y_train)
```

```
[ ]: # pram_grid = {
#     'logisticregression__C': [0.01, 0.1, 1, 10], # Assuming 'A' was intended
#     'columntransformer__pipeline-1__simpleimputer__strategy':
#     'mean', 'median', # Strategy for SimpleImputer in age_pipe2
#     'columntransformer__pipeline-3__simpleimputer__strategy':
#     'most_frequent', 'constant' # Strategy for SimpleImputer in sex_emb_pipe2
# }

# grid_search=GridSearchCV(pipe3, pram_grid, cv=5, scoring='accuracy')
# grid_search.fit(X_train,y_train)
```

#Day-39:KNN Imputer

```
[ ]: X_train.head()
```

```
[ ]:
```

	Pclass	Sex	Age	SibSp	Parch	Fare	Embarked	new_age
331	1	male	45.5	0	0	28.5000	S	False
733	2	male	23.0	0	0	13.0000	S	False
382	3	male	32.0	0	0	7.9250	S	False

704	3	male	26.0	1	0	7.8542	S	False
813	3	female	6.0	4	2	31.2750	S	False

```
[ ]:
```

```
[ ]: X_train['Age'].isnull().sum()
```

```
[ ]: np.int64(140)
```

```
[ ]: #ChatGPT
```

```
from sklearn.pipeline import Pipeline
from sklearn.impute import KNNImputer
from sklearn.preprocessing import PowerTransformer
from sklearn.compose import ColumnTransformer
from sklearn.linear_model import LogisticRegression
from sklearn.preprocessing import OneHotEncoder
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score
import pandas as pd

# 1. Load Titanic dataset
df = df
X = df.drop('Survived', axis=1)
y = df['Survived']

# 2. Split data
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    random_state=42)

# 3. Select columns
numeric_cols = ['Age', 'Fare', 'SibSp', 'Parch']
categorical_cols = ['Sex', 'Embarked']

# 4. Step 1: Numeric imputation using KNN
knn_impute_pipe = Pipeline([
    ('knn_imputer', KNNImputer())
])

# 5. Step 2: PowerTransform only Fare
fare_power_pipe = Pipeline([
    ('power', PowerTransformer(method='yeo-johnson'))
])

# 6. Step 3: Categorical encode
cat_pipe = Pipeline([
```

```

        ('onehot', OneHotEncoder(handle_unknown='ignore'))
    ])

# 7. Combine: ColumnTransformer
preprocessor = ColumnTransformer([
    ('knn_num', knn_impute_pipe, numeric_cols), # shob golo column use kore just
    ↪akta column "Age" ei value Add korse[Onno kono Column k Affact kore nai. Tai
    ↪kaj korche]
    #Ta o jodi kaj na kore Tahole
    ↪Fare Column k bad dia all Numeric column pass koro
    ('fare_power', fare_power_pipe, ['Fare']),
    ('cat', cat_pipe, categorical_cols)
])

# 8. Final Pipeline: Preprocessor + Model
model_pipeline = Pipeline([
    ('preprocess', preprocessor),
    ('clf', LogisticRegression(max_iter=1000))
])

# 9. Train
model_pipeline.fit(X_train, y_train)

# 10. Predict & Evaluate
y_pred = model_pipeline.predict(X_test)
print("Accuracy:", accuracy_score(y_test, y_pred))

```

Accuracy: 0.770949720670391

```

[ ]: from sklearn.impute import KNNImputer
knn_age_pipe=Pipeline([
    ('knn_age',KNNImputer()) #just Age colum Pass korle hobe na All Numerical
    ↪Column Pass koro(with Encoded if have)
])

```

```

[ ]: CT4=ColumnTransformer([
    ('knn_age_imp',knn_age_pipe,['Age']),#just Age ar opor KNN apply kore Fill
    ↪kore.But KNN All column ar distance ar opor bitte kore value ber korar kotha
    ('normalize_fare',fare_pipe,['Fare']),
    ('sex_emb_ohe',sex_emb_pipe,['Sex','Embarked']),
    ],remainder='passthrough')

```

```

[ ]: pipe_knn=make_pipeline(
    CT4,StandardScaler(),LogisticRegression()
)
pipe_knn.fit(X_train,y_train)

```

/usr/local/lib/python3.11/dist-packages/sklearn/compose/_column_transformer.py:1667: FutureWarning:
The format of the columns of the 'remainder' transformer in
ColumnTransformer.transformers_ will change in version 1.7 to match the format
of the other transformers.
At the moment the remainder columns are stored as indices (of type int). With
the same ColumnTransformer configuration, in the future they will be stored as
column names (of type str).
To use the new behavior now and suppress this warning, use
ColumnTransformer(force_int_remainder_cols=False).

```
warnings.warn(
```

```
[ ]: Pipeline(steps=[('columntransformer',
                      ColumnTransformer(remainder='passthrough',
                      transformers=[('knn_age_imp',
                                   Pipeline(steps=[('knn_age',
                                                    Pipeline(steps=[('knn_age',
                                                                    ['Age'])),
                                                                    ('normalize_fare',
                                                                    ['Fare'])),
                                                                    ('sex_emb_ohe',
                                                                    Pipeline(steps=[('sex_imp',
                                                                    SimpleImputer(strategy='most_frequent')),
                                                                    ('ohe_sex_emb',
                                                                    OneHotEncoder(drop='first',
                                                                    dtype=<class 'numpy.int32'>,
                                                                    handle_unknown='ignore',
                                                                    sparse_output=False)))]),
                                                                    ['Sex', 'Embarked']]])),
                      ('standardscaler', StandardScaler()),
                      ('logisticregression', LogisticRegression())])
```

```
[ ]: y_pred_knn=pipe_knn.predict(X_test)
      accuracy_score(y_pred_knn,y_test)
```

```
[ ]: 0.8100558659217877
```

```
[ ]: cross_val_score(pipe_knn,X_train,y_train,cv=5,scoring='accuracy').mean()
```

```
[ ]: np.float64(0.7864572047670639)
```

```
[ ]:
```

```
#IterativeImputer:(use Linear Model)
```

Iterative Imputer

```
###          ? IterativeImputer
            multivariate imputation
```

Missing value handling using Iterative Imputer (from sklearn.impute)

0.0.1 ?

IterativeImputer
multivariate imputation

0.0.2 : Python (Scikit-learn)

```
import pandas as pd
import numpy as np
from sklearn.impute import IterativeImputer
from sklearn.linear_model import BayesianRidge #

#
data = {
    'age': [25, np.nan, 35, 40, np.nan],
    'salary': [50000, 60000, np.nan, 80000, 75000],
    'experience': [2, 4, 5, np.nan, 3]
}
df = pd.DataFrame(data)

#
imputer = IterativeImputer(estimator=BayesianRidge(), max_iter=10, random_state=0)

#
df_imputed = imputer.fit_transform(df)

#      DataFrame
df_imputed = pd.DataFrame(df_imputed, columns=df.columns)

print(df_imputed)
```

0.0.3 Parameters

- estimator: (BayesianRidge, DecisionTreeRegressor, KNN)
- max_iter:
- initial_strategy: (mean, median, most_frequent)

- `random_state` :Reproducibility `random_state`
-

0.0.4 ?

-
- or
- Simple mean/median
-

0.1 (SimpleImputer)

0.1.1 :

- (slow)
- (StandardScaler)
- High Memory Consumed in for Server(deploy with main dataset for user missing value handling.
- Overfitting: Iteration

```
[ ]: from sklearn.pipeline import Pipeline
from sklearn.experimental import enable_iterative_imputer #
from sklearn.impute import IterativeImputer
from sklearn.preprocessing import PowerTransformer
from sklearn.compose import ColumnTransformer
from sklearn.linear_model import LogisticRegression
from sklearn.preprocessing import OneHotEncoder
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score
import pandas as pd

# 1. Load Titanic dataset
df = df
X = df.drop('Survived', axis=1)
y = df['Survived']

# 2. Split data
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳random_state=42)

# 3. Select columns
numeric_cols = ['Age', 'Fare', 'SibSp', 'Parch']
categorical_cols = ['Sex', 'Embarked']

# 4. Step 1: Numeric imputation using IterativeImputer
iter_impute_pipe = Pipeline([
    ('iter_imputer', IterativeImputer(random_state=42))
```

```

])

# 5. Step 2: PowerTransform only Fare
fare_power_pipe = Pipeline([
    ('power', PowerTransformer(method='yeo-johnson'))
])

# 6. Step 3: Categorical encode
cat_pipe = Pipeline([
    ('onehot', OneHotEncoder(handle_unknown='ignore'))
])

# 7. Combine: ColumnTransformer
preprocessor = ColumnTransformer([
    ('iter_num', iter_impute_pipe, numeric_cols),
    ('fare_power', fare_power_pipe, ['Fare']),
    ('cat', cat_pipe, categorical_cols)
])

# 8. Final Pipeline: Preprocessor + Model
model_pipeline = Pipeline([
    ('preprocess', preprocessor),
    ('clf', LogisticRegression(max_iter=1000))
])

# 9. Train
model_pipeline.fit(X_train, y_train)

# 10. Predict & Evaluate
y_pred = model_pipeline.predict(X_test)
print("Accuracy with Iterative Imputer:", accuracy_score(y_test, y_pred))

```

Accuracy with Iterative Imputer: 0.770949720670391

```

[ ]: from sklearn.pipeline import Pipeline
from sklearn.experimental import enable_iterative_imputer
from sklearn.impute import IterativeImputer
from sklearn.preprocessing import PowerTransformer, OneHotEncoder
from sklearn.compose import ColumnTransformer
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score
import pandas as pd

# Load Titanic dataset
# Make sure df is already loaded properly with 'Survived' column
X = df.drop('Survived', axis=1)

```

```

y = df['Survived']

# Split data
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    random_state=42)

# Column groups
numeric_cols = ['Age', 'SibSp', 'Parch']          # Only these go to
    IterativeImputer
fare_col = ['Fare']                                # Separate transformer
categorical_cols = ['Sex', 'Embarked']

# Pipelines
iter_impute_pipe = Pipeline([
    ('iter_imputer', IterativeImputer(estimator=LogisticRegression(),
    random_state=42))
])

fare_power_pipe = Pipeline([
    ('power', PowerTransformer(method='yeo-johnson'))
])

cat_pipe = Pipeline([
    ('onehot', OneHotEncoder(handle_unknown='ignore'))
])

# ColumnTransformer
preprocessor = ColumnTransformer([
    ('iter_num', iter_impute_pipe, numeric_cols),
    ('fare_power', fare_power_pipe, fare_col),
    ('cat', cat_pipe, categorical_cols)
])

# Final pipeline
model_pipeline = Pipeline([
    ('preprocess', preprocessor),
    ('clf', LogisticRegression(max_iter=1000))
])

# Train
model_pipeline.fit(X_train, y_train)

# Predict & Evaluate
y_pred = model_pipeline.predict(X_test)
print("Accuracy with Iterative Imputer:", accuracy_score(y_test, y_pred))

```

/usr/local/lib/python3.11/dist-packages/sklearn/linear_model/_logistic.py:465:


```
ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. OF ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
/usr/local/lib/python3.11/dist-packages/sklearn/linear_model/_logistic.py:465:
ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. OF ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

```
-----
ValueError                                Traceback (most recent call last)
<ipython-input-126-00f56afa5758> in <cell line: 0>()
    49
    50 # Train
--> 51 model_pipeline.fit(X_train, y_train)
    52
    53 # Predict & Evaluate

/usr/local/lib/python3.11/dist-packages/sklearn/base.py in wrapper(estimator, *
↳ *args, **kwargs)
    1387         )
    1388     ):
-> 1389         return fit_method(estimator, *args, **kwargs)
    1390
    1391     return wrapper

/usr/local/lib/python3.11/dist-packages/sklearn/pipeline.py in fit(self, X, y,
↳ **params)
    652
    653     routed_params = self._check_method_params(method="fit",
↳ props=params)
--> 654     Xt = self._fit(X, y, routed_params, raw_params=params)
    655     with _print_elapsed_time("Pipeline", self._log_message(len(self
↳ steps) - 1)):
    656         if self._final_estimator != "passthrough":
```

```

/usr/local/lib/python3.11/dist-packages/sklearn/pipeline.py in _fit(self, X, y,
↳ routed_params, raw_params)
    586         )
    587
--> 588         X, fitted_transformer = fit_transform_one_cached(
    589             cloned_transformer,
    590             X,

/usr/local/lib/python3.11/dist-packages/joblib/memory.py in __call__(self,
↳ *args, **kwargs)
    324
    325     def __call__(self, *args, **kwargs):
--> 326         return self.func(*args, **kwargs)
    327
    328     def call_and_shelve(self, *args, **kwargs):

/usr/local/lib/python3.11/dist-packages/sklearn/pipeline.py in
↳ _fit_transform_one(transformer, X, y, weight, message_clsname, message, param
    1549         with _print_elapsed_time(message_clsname, message):
    1550             if hasattr(transformer, "fit_transform"):
-> 1551                 res = transformer.fit_transform(X, y, **params.
↳ get("fit_transform", {}))
    1552             else:
    1553                 res = transformer.fit(X, y, **params.get("fit", {})).
↳ transform(

/usr/local/lib/python3.11/dist-packages/sklearn/utils/_set_output.py in
↳ wrapped(self, X, *args, **kwargs)
    317     @wraps(f)
    318     def wrapped(self, X, *args, **kwargs):
--> 319         data_to_wrap = f(self, X, *args, **kwargs)
    320         if isinstance(data_to_wrap, tuple):
    321             # only wrap the first output for cross decomposition

/usr/local/lib/python3.11/dist-packages/sklearn/base.py in wrapper(estimator,
↳ *args, **kwargs)
    1387         )
    1388     ):
-> 1389         return fit_method(estimator, *args, **kwargs)
    1390
    1391     return wrapper

/usr/local/lib/python3.11/dist-packages/sklearn/compose/_column_transformer.py
↳ in fit_transform(self, X, y, **params)
    999         routed_params = self._get_empty_routing()
    1000

```

```

-> 1001         result = self._call_func_on_transformers(
    1002             X,
    1003             y,

/usr/local/lib/python3.11/dist-packages/sklearn/compose/_column_transformer.py
↳in _call_func_on_transformers(self, X, y, func, column_as_labels,
↳routed_params)
    908         )
    909
-> 910         return Parallel(n_jobs=self.n_jobs)(jobs)
    911
    912     except ValueError as e:

/usr/local/lib/python3.11/dist-packages/sklearn/utils/parallel.py in
↳__call__(self, iterable)
    75         for delayed_func, args, kwargs in iterable
    76     )
---> 77     return super().__call__(iterable_with_config)
    78
    79

/usr/local/lib/python3.11/dist-packages/joblib/parallel.py in __call__(self,
↳iterable)
    1983         output = self._get_sequential_output(iterable)
    1984         next(output)
-> 1985         return output if self.return_generator else list(output)
    1986
    1987         # Let's create an ID that uniquely identifies the current call.
↳If the

/usr/local/lib/python3.11/dist-packages/joblib/parallel.py in
↳_get_sequential_output(self, iterable)
    1911         self.n_dispatched_batches += 1
    1912         self.n_dispatched_tasks += 1
-> 1913         res = func(*args, **kwargs)
    1914         self.n_completed_tasks += 1
    1915         self.print_progress()

/usr/local/lib/python3.11/dist-packages/sklearn/utils/parallel.py in
↳__call__(self, *args, **kwargs)
    137         config = {}
    138         with config_context(**config):
-> 139         return self.function(*args, **kwargs)
    140
    141

```

```

/usr/local/lib/python3.11/dist-packages/sklearn/pipeline.py in
-> _fit_transform_one(transformer, X, y, weight, message_clsname, message, param
1549     with _print_elapsed_time(message_clsname, message):
1550         if hasattr(transformer, "fit_transform"):
-> 1551             res = transformer.fit_transform(X, y, **params.
-> get("fit_transform", {}))
1552         else:
1553             res = transformer.fit(X, y, **params.get("fit", {})).
-> transform(

/usr/local/lib/python3.11/dist-packages/sklearn/base.py in wrapper(estimator,
-> *args, **kwargs)
1387         )
1388     ):
-> 1389         return fit_method(estimator, *args, **kwargs)
1390
1391     return wrapper

/usr/local/lib/python3.11/dist-packages/sklearn/pipeline.py in
-> fit_transform(self, X, y, **params)
728         )
729         if hasattr(last_step, "fit_transform"):
--> 730             return last_step.fit_transform(
731                 Xt, y, **last_step_params["fit_transform"]
732             )

/usr/local/lib/python3.11/dist-packages/sklearn/utils/_set_output.py in
-> wrapped(self, X, *args, **kwargs)
317     @wraps(f)
318     def wrapped(self, X, *args, **kwargs):
--> 319         data_to_wrap = f(self, X, *args, **kwargs)
320         if isinstance(data_to_wrap, tuple):
321             # only wrap the first output for cross decomposition

/usr/local/lib/python3.11/dist-packages/sklearn/base.py in wrapper(estimator,
-> *args, **kwargs)
1387         )
1388     ):
-> 1389         return fit_method(estimator, *args, **kwargs)
1390
1391     return wrapper

/usr/local/lib/python3.11/dist-packages/sklearn/impute/_iterative.py in
-> fit_transform(self, X, y, **params)
857         n_features, feat_idx, abs_corr_mat
858     )

```

```

--> 859             Xt, estimator = self._impute_one_feature(
            860                 Xt,
            861                 mask_missing_values,

/usr/local/lib/python3.11/dist-packages/sklearn/impute/_iterative.py in
↳ _impute_one_feature(self, X_filled, mask_missing_values, feat_idx,
↳ neighbor_feat_idx, estimator, fit_mode, params)
            425                 axis=0,
            426             )
--> 427             estimator.fit(X_train, y_train, **params)
            428
            429             # if no missing values, don't predict

/usr/local/lib/python3.11/dist-packages/sklearn/base.py in wrapper(estimator,
↳ *args, **kwargs)
            1387         )
            1388     ):
-> 1389         return fit_method(estimator, *args, **kwargs)
            1390
            1391     return wrapper

/usr/local/lib/python3.11/dist-packages/sklearn/linear_model/_logistic.py in
↳ fit(self, X, y, sample_weight)
            1229         accept_large_sparse=solver not in ["liblinear", "sag",
↳ "saga"],
            1230     )
-> 1231     check_classification_targets(y)
            1232     self.classes_ = np.unique(y)
            1233

/usr/local/lib/python3.11/dist-packages/sklearn/utils/multiclass.py in
↳ check_classification_targets(y)
            220         "multilabel-sequences",
            221     ]:
--> 222         raise ValueError(

            223             f"Unknown label type: {y_type}. Maybe you are trying to fit
↳ a "
            224             "classifier, which expects discrete classes on a "

ValueError: Unknown label type: continuous. Maybe you are trying to fit a
↳ classifier, which expects discrete classes on a regression target with
↳ continuous values.

```